

## Machine Learning Lab

```
# NAME :          Bhuvnesh Verma
# ROLL NO :       28
# SECTION :       A
# LAB NO. :       03
# OBJECTIVE :     Data Preprocessing and Survival Prediction
```

```
'''
```

```
Aim : To predict Titanic survivors using a Random Forest Classifier by preprocessing data,
engineering features, and applying PCA for dimensionality reduction.
'''
```

```
# Load the Titanic dataset from a URL
url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"
data = pd.read_csv(url)
```

```
# Display the first few rows of the dataset to understand its structure
print(data.head())
```

```
➡ PassengerId  Survived  Pclass \
0            1         0       3
1            2         1       1
2            3         1       3
3            4         1       1
4            5         0       3

      Name      Sex  Age  SibSp \
0   Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2     Heikkinen, Miss. Laina  female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4     Allen, Mr. William Henry    male  35.0      0

      Parch      Ticket    Fare Cabin Embarked
0         0   A/5 21171   7.2500   NaN        S
1         0    PC 17599  71.2833   C85        C
2         0  STON/O2. 3101282   7.9250   NaN        S
3         0    113803   53.1000  C123        S
4         0   373450   8.0500   NaN        S
```

```
# Displays dataframe info and data types
data.info()
```

```
➡ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column        Non-Null Count  Dtype  
---  -
0   PassengerId    891 non-null   int64  
1   Survived       891 non-null   int64  
2   Pclass        891 non-null   int64  
3   Name          891 non-null   object  
4   Sex           891 non-null   object  
5   Age           714 non-null   float64 
6   SibSp         891 non-null   int64  
7   Parch         891 non-null   int64  
8   Ticket        891 non-null   object  
9   Fare          891 non-null   float64 
10  Cabin         204 non-null   object  
11  Embarked      889 non-null   object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
# Show summary statistics of numeric columns
data.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
# Check for missing values in the dataset
print(data.isnull().sum())
```

```
PassengerId    0
Survived       0
Pclass         0
Name           0
Sex            0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64
```

```
# Drop unnecessary columns that are not relevant for the analysis
data = data.drop(columns=["PassengerId", "Name", "Ticket", "Cabin", "Fare"])
```

```
# Display the first few rows after dropping columns
print(data.head())
```

```
Survived  Pclass  Sex  Age  SibSp  Parch  Embarked
0         0      3  male  22.0    1      0         S
1         1      1  female  38.0    1      0         C
2         1      3  female  26.0    0      0         S
3         1      1  female  35.0    1      0         S
4         0      3  male   35.0    0      0         S
```

```
# As there are missing value in Age cloumn
# Fill missing values in the 'Age' column with the median value
data["Age"] = data["Age"].fillna(data["Age"].median())
```

```
# As there are missing value in Embarked column. Use mode as we have only S & C
# Fill missing values in the 'Embarked' column with the mode value
data["Embarked"] = data["Embarked"].fillna(data["Embarked"].mode()[0])
```

```
# Display the first few rows after filling missing values
print(data.head())
```

```
Survived  Pclass  Sex  Age  SibSp  Parch  Embarked
0         0      3  male  22.0    1      0         S
1         1      1  female  38.0    1      0         C
2         1      3  female  26.0    0      0         S
3         1      1  female  35.0    1      0         S
4         0      3  male   35.0    0      0         S
```

```
# SibSp - the number of siblings and spouses a passenger had on board.
# Parch - the number of parents and children a passenger had aboard the ship.
# Create a new feature 'family_size' by combining 'SibSp' and 'Parch' columns
data["family_size"] = data["SibSp"] + data["Parch"] + 1
```

```
# Create a new feature 'Is_Alone' to indicate if a passenger is traveling alone
data["Is_Alone"] = (data["family_size"] == 1).astype(int)
```

```
# Embarked - indicates the port of embarkation for each passenger. [C: Cherbourg , Q: Queenstown , S: Southampton]
# Create a new feature 'Title' by combining 'Sex' and 'Embarked' columns
data["Title"] = data["Sex"] + "_" + data["Embarked"]
```

```
# Display the first few rows after creating new features
print(data.head())
```

```

Survived  Pclass    Sex  Age  SibSp  Parch  Embarked  family_size \
0         0        3  male  22.0    1     0         S           2
1         1        1 female  38.0    1     0         C           2
2         1        3 female  26.0    0     0         S           1
3         1        1 female  35.0    1     0         S           2
4         0        3  male   35.0    0     0         S           1

Is_Alone  Title
0         0  male_S
1         0 female_C
2         1 female_S
3         0 female_S
4         1  male_S

```

```
# Convert categorical columns into binary columns using one-hot
...
```

```
Convert to binary features called dummy variables
```

```
0: Observation was NOT that category
```

```
1: Observation was that category
```

```
...
```

```
data = pd.get_dummies(data, columns=["Sex", "Embarked", "Title"], drop_first=True)
```

```
# Drop 'SibSp' and 'Parch' columns as 'family_size' is a better representation
```

```
data = data.drop(columns=["SibSp", "Parch"])
```

```
# Display the first few rows after removing SibSp & Parch
```

```
print(data.head())
```

```

Survived  Pclass  Age  family_size  Is_Alone  Sex_male  Embarked_Q \
0         0        3  22.0           2         0      True      False
1         1        1  38.0           2         0     False      False
2         1        3  26.0           1         1     False      False
3         1        1  35.0           2         0     False      False
4         0        3  35.0           1         1      True      False

Embarked_S  Title_female_Q  Title_female_S  Title_male_C  Title_male_Q \
0         True           False           False           False           False
1        False           False           False           False           False
2         True           False           True           False           False
3         True           False           True           False           False
4         True           False           False           False           False

Title_male_S
0         True
1        False
2        False
3        False
4         True

```

```
# Split the dataset into features (X) and target (y)
```

```
X = data.drop(columns=["Survived"])
```

```
y = data["Survived"]
```

```
# Standardize the features to have a mean of 0 and a standard deviation of 1
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# Split the dataset into training and testing sets (80% training, 20% testing)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Apply Principal Component Analysis (PCA) to reduce dimensionality
```

```
pca = PCA(n_components=12)
```

```
X_train_pca = pca.fit_transform(X_train)
```

```
X_test_pca = pca.fit_transform(X_test)
```

```
# Print the explained variance ratio of each principal component
```

```
print("Explained variance ratio", pca.explained_variance_ratio_)
```

```

Explained variance ratio [9.79021339e-01 1.04195229e-02 4.17053134e-03 2.77165865e-03
 1.83105681e-03 7.66358339e-04 4.80432941e-04 3.02982410e-04
 2.36117755e-04 3.48823904e-19 2.54973412e-20 0.00000000e+00]

```

```
# Initialize a Random Forest Classifier with a fixed random state for reproducibility
```

```
rf = RandomForestClassifier(random_state=42)
```

```
# Train the Random Forest model on the training data
rf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
# Predict the target values for the test set
y_pred = rf.predict(X_test)
```

```
# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy :", accuracy)
```

```
Accuracy : 0.8100558659217877
```

```
# Calculate the Mean Squared Error (MSE) and R² Score to evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error : {mse}")
print(f"R² Score : {r2}")
```

```
Mean Squared Error : 0.18994413407821228
R² Score : 0.2167310167310167
```

Here is a concise explanation of **PCA** and **Random Forest Classifier**

## 1. Principal Component Analysis (PCA)

### What is PCA?

- A dimensionality reduction technique.
- Transforms original features into new uncorrelated features called **principal components**.

### Why Use PCA?

- Reduces dataset complexity.
- Removes redundant or correlated features.
- Improves visualization of high-dimensional data.
- Speeds up machine learning algorithms.

### How PCA Works

1. **Standardize Data:** Scale features to have mean = 0 and standard deviation = 1.
2. **Compute Covariance Matrix:** Measures how features vary together.
3. **Calculate Eigenvalues and Eigenvectors:** Eigenvectors = directions (principal components), eigenvalues = variance along those directions.
4. **Sort and Select Components:** Keep top *n* components with the highest variance.
5. **Transform Data:** Project data onto the selected components.

### Key Terms

- **Principal Components:** New uncorrelated features.
- **Explained Variance Ratio:** Proportion of variance captured by each component.
- **Dimensionality Reduction:** Reducing the number of features while retaining information.

### Example Code

```
from sklearn.decomposition import PCA
pca = PCA(n_components=12)
X_train_pca = pca.fit_transform(X_train)
print("Explained variance ratio:", pca.explained_variance_ratio_)
```

---

## 2. Random Forest Classifier

### What is Random Forest?

- An ensemble learning method for classification and regression.
- Combines multiple decision trees to improve accuracy and reduce overfitting.

Why Use Random Forest?

- High accuracy due to ensemble averaging.
- Robust to overfitting.
- Handles missing data and outliers.
- Provides feature importance scores.

How Random Forest Works

1. **Bootstrap Sampling:** Create random subsets of the training data for each tree.
2. **Feature Randomness:** At each split, consider a random subset of features.
3. **Build Decision Trees:** Train multiple trees on different subsets.
4. **Voting/Averaging:** For classification, use majority voting; for regression, use averaging.

Key Terms

- **Ensemble Learning:** Combining multiple models for better performance.
- **Decision Trees:** Individual models that make predictions based on if-else conditions.
- **Feature Importance:** Scores indicating the contribution of each feature to predictions.

Example Code

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Comparison: PCA vs Random Forest

Aspect	PCA	Random Forest
Purpose	Dimensionality reduction	Classification/Regression
Output	Transformed features (principal components)	Predictions (e.g., class labels or continuous values)
Key Idea	Reduce features while retaining variance	Combine multiple decision trees to improve accuracy
Use Case	Preprocessing step for high-dimensional data	Final model for making predictions
Interpretability	Principal components are not directly interpretable	Provides feature importance for interpretability
Handles Overfitting	Reduces overfitting by reducing dimensionality	Reduces overfitting through ensemble averaging

When to Use PCA and Random Forest Together?

- **High-Dimensional Data:** Use PCA to reduce features before training Random Forest.
- **Improving Performance:** PCA removes noise, which can improve Random Forest's accuracy.
- **Visualization:** PCA helps visualize data in 2D/3D, while Random Forest makes predictions.

Workflow of Practical

1. **Load Data:** Load dataset (e.g., Titanic dataset).
2. **Preprocess Data:** Handle missing values, encode categorical variables, and create new features.
3. **Apply PCA:** Reduce dimensionality using PCA.
4. **Train Random Forest:** Train the model on the reduced dataset.
5. **Evaluate Model:** Check accuracy, MSE, and R² score.

```
...
Conclusion: The model achieved good accuracy,
highlighting the importance of feature selection,
data cleaning, and dimensionality reduction in
improving prediction performance. 📊💡
...
```